


End-to-end task latency per run



Dots: individual runs. Bars: median. Same functions, images and mock LLM on both stacks.